

```
In [ ]: !pip install ISLP
!pip install catboost
```

Collecting ISLP

Downloading ISLP-0.4.0-py3-none-any.whl.metadata (7.0 kB)
Requirement already satisfied: numpy>=1.7.1 in /usr/local/lib/python3.10/dist-packages (from ISLP) (1.26.4)
Requirement already satisfied: scipy>=0.9 in /usr/local/lib/python3.10/dist-packages (from ISLP) (1.13.1)
Requirement already satisfied: pandas>=0.20 in /usr/local/lib/python3.10/dist-packages (from ISLP) (2.2.2)
Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from ISLP) (5.3.0)
Requirement already satisfied: scikit-learn>=1.2 in /usr/local/lib/python3.10/dist-packages (from ISLP) (1.5.2)
Requirement already satisfied: joblib in /usr/local/lib/python3.10/dist-packages (from ISLP) (1.4.2)
Requirement already satisfied: statsmodels>=0.13 in /usr/local/lib/python3.10/dist-packages (from ISLP) (0.14.4)
Collecting lifelines (from ISLP)
Downloading lifelines-0.30.0-py3-none-any.whl.metadata (3.2 kB)
Collecting pygam (from ISLP)
Downloading pygam-0.9.1-py3-none-any.whl.metadata (7.1 kB)
Requirement already satisfied: torch in /usr/local/lib/python3.10/dist-packages (from ISLP) (2.5.1+cu121)
Collecting pytorch-lightning (from ISLP)
Downloading pytorch_lightning-2.4.0-py3-none-any.whl.metadata (21 kB)
Collecting torchmetrics (from ISLP)
Downloading torchmetrics-1.6.0-py3-none-any.whl.metadata (20 kB)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20->ISLP) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20->ISLP) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.20->ISLP) (2024.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.2->ISLP) (3.5.0)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels>=0.13->ISLP) (1.0.1)
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.10/dist-packages (from statsmodels>=0.13->ISLP) (24.2)
Requirement already satisfied: matplotlib>=3.0 in /usr/local/lib/python3.10/dist-packages (from lifelines->ISLP) (3.8.0)
Requirement already satisfied: autograd>=1.5 in /usr/local/lib/python3.10/dist-packages (from lifelines->ISLP) (1.7.0)
Collecting autograd-gamma>=0.3 (from lifelines->ISLP)
Downloading autograd-gamma-0.5.0.tar.gz (4.0 kB)
Preparing metadata (setup.py) ... done
Collecting formulaic>=0.2.2 (from lifelines->ISLP)
Downloading formulaic-1.0.2-py3-none-any.whl.metadata (6.8 kB)
Requirement already satisfied: progressbar2<5.0.0,>=4.2.0 in /usr/local/lib/python3.10/dist-packages (from pygam->ISLP) (4.5.0)
Collecting scipy>=0.9 (from ISLP)
Downloading scipy-1.11.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl.metadata (60 kB)

60.4/60.4 kB 3.8 MB/s eta 0:00:00
Requirement already satisfied: tqdm>=4.57.0 in /usr/local/lib/python3.10/dist-packages (from pytorch-lightning->ISLP) (4.66.6)
Requirement already satisfied: PyYAML>=5.4 in /usr/local/lib/python3.10/dist-packages (from pytorch-lightning->ISLP) (6.0.2)
Requirement already satisfied: fsspec>=2022.5.0 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (2024.10.0)
Requirement already satisfied: typing-extensions>=4.4.0 in /usr/local/lib/python3.10/dist-packages (from pytorch-lightning->ISLP) (4.12.2)
Collecting lightning-utilities>=0.10.0 (from pytorch-lightning->ISLP)
Downloading lightning_utilities-0.11.9-py3-none-any.whl.metadata (5.2 kB)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch->ISLP) (3.16.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch->ISLP) (3.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from torch->ISLP) (3.1.4)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.10/dist-packages (from torch->ISLP) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from sympy==1.13.1->torch->ISLP) (1.3.0)
Collecting interface-meta>=1.2.0 (from formulaic>=0.2.2->lifelines->ISLP)
Downloading interface_meta-1.3.0-py3-none-any.whl.metadata (6.7 kB)
Requirement already satisfied: wrapt>=1.0 in /usr/local/lib/python3.10/dist-packages (from formulaic>=0.2.2->lifelines->ISLP) (1.17.0)
Requirement already satisfied: aiohttp!=4.0.0a0,!4.0.0a1 in /usr/local/lib/python3.10/dist-packages (from fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (3.11.9)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from lightning-utilities>=0.10.0->pytorch-lightning->ISLP) (75.1.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0->lifelines->ISLP) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0->lifelines->ISLP) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0->lifelines->ISLP) (4.55.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0->lifelines->ISLP) (1.4.7)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0->lifelines->ISLP) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.0->lifelines->ISLP) (3.2.0)
Requirement already satisfied: python-utils>=3.8.1 in /usr/local/lib/python3.10/dist-packages (from progressbar2<5.0.0,>=4.2.0->pygam->ISLP) (3.9.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.8.2->pandas>=0.20->ISLP) (1.16.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from Jinja2->torch->ISLP) (3.0.2)
Requirement already satisfied: aiohappyeyeballs>=2.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (2.4.4)
Requirement already satisfied: aiosignal>=1.1.2 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (1.3.1)
Requirement already satisfied: async-timeout<6.0,>=4.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (4.0.3)
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (24.2.0)
Requirement already satisfied: frozenlist>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1->fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (1.5.0)
Requirement already satisfied: multidict<7.0,>=4.5 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!4.0.0a1-

```

>fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (6.1.0)
Requirement already satisfied: propcache>=0.2.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!
spec[http]>=2022.5.0->pytorch-lightning->ISLP) (0.2.1)
Requirement already satisfied: yarl<2.0,>=1.17.0 in /usr/local/lib/python3.10/dist-packages (from aiohttp!=4.0.0a0,!
sspec[http]>=2022.5.0->pytorch-lightning->ISLP) (1.18.3)
Requirement already satisfied: idna>=2.0 in /usr/local/lib/python3.10/dist-packages (from yarl<2.0,>=1.17.0->aiohttp!=4.0.0a0,!
=4.0.0a1->fsspec[http]>=2022.5.0->pytorch-lightning->ISLP) (3.10)
Downloading ISLP-0.4.0-py3-none-any.whl (3.6 MB)
_____ 3.6/3.6 MB 52.7 MB/s eta 0:00:00
Downloading lifelines-0.30.0-py3-none-any.whl (349 kB)
_____ 349.3/349.3 kB 22.4 MB/s eta 0:00:00
Downloading pygam-0.9.1-py3-none-any.whl (522 kB)
_____ 522.0/522.0 kB 29.7 MB/s eta 0:00:00
Downloading scipy-1.11.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (36.4 MB)
_____ 36.4/36.4 MB 21.5 MB/s eta 0:00:00
Downloading pytorch_lightning-2.4.0-py3-none-any.whl (815 kB)
_____ 815.2/815.2 kB 10.0 MB/s eta 0:00:00
Downloading torchmetrics-1.6.0-py3-none-any.whl (926 kB)
_____ 926.4/926.4 kB 46.1 MB/s eta 0:00:00
Downloading formulaic-1.0.2-py3-none-any.whl (94 kB)
_____ 94.5/94.5 kB 7.0 MB/s eta 0:00:00
Downloading lightning_utilities-0.11.9-py3-none-any.whl (28 kB)
Downloading interface_meta-1.3.0-py3-none-any.whl (14 kB)
Building wheels for collected packages: autograd-gamma
  Building wheel for autograd-gamma (setup.py) ... done
  Created wheel for autograd-gamma: filename=autograd_gamma-0.5.0-py3-none-any.whl size=4031 sha256=4f9e7490640d47feffb5257bb51
58ba92c129d664024a22c9b3f84708c5784dc
  Stored in directory: /root/.cache/pip/wheels/25/cc/e0/ef2969164144c899fedb22b338f6703e2b9cf46eeebf254991
Successfully built autograd-gamma
Installing collected packages: scipy, lightning-utilities, interface-meta, autograd-gamma, torchmetrics, pygam, formulaic, life
lines, pytorch-lightning, ISLP
  Attempting uninstall: scipy
    Found existing installation: scipy 1.13.1
    Uninstalling scipy-1.13.1:
      Successfully uninstalled scipy-1.13.1
Successfully installed ISLP-0.4.0 autograd-gamma-0.5.0 formulaic-1.0.2 interface-meta-1.3.0 lifelines-0.30.0 lightning-utilitie
s-0.11.9 pygam-0.9.1 pytorch-lightning-2.4.0 scipy-1.11.4 torchmetrics-1.6.0
Collecting catboost
  Downloading catboost-1.2.7-cp310-cp310-manylinux2014_x86_64.whl.metadata (1.2 kB)
Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (from catboost) (0.20.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from catboost) (3.8.0)
Requirement already satisfied: numpy<2.0,>=1.16.0 in /usr/local/lib/python3.10/dist-packages (from catboost) (1.26.4)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.10/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packages (from catboost) (1.11.4)
Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.10/dist-packages (from catboost) (1.16.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost)
(2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2024.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas>=0.24->catboost) (2024.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.3.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (4.55.
1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (1.4.7)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (24.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (11.0.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->catboost) (3.2.0)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly->catboost) (9.0.0)
Downloading catboost-1.2.7-cp310-cp310-manylinux2014_x86_64.whl (98.7 MB)
_____ 98.7/98.7 MB 7.7 MB/s eta 0:00:00
Installing collected packages: catboost
Successfully installed catboost-1.2.7

```

```

In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib.pyplot import subplots
import statsmodels.api as sm
from ISLP import load_data
from ISLP.models import (ModelSpec as MS,
summarize)

from ISLP import confusion_table
from ISLP.models import contrast
from sklearn.discriminant_analysis import \
(LinearDiscriminantAnalysis as LDA,
QuadraticDiscriminantAnalysis as QDA)
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from itertools import combinations

import seaborn as sns
from sklearn.tree import DecisionTreeClassifier

```


```
from sklearn.metrics import accuracy_score, f1_score, recall_score, balanced_accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import GridSearchCV
from imblearn.over_sampling import SMOTE
from sklearn.ensemble import RandomForestClassifier
import xgboost as xgb
from sklearn.linear_model import LogisticRegression
```

```
In [ ]: df = pd.read_csv("train.csv")
df.head(5)
```

Out[]:

	PatientID	Age	Gender	Ethnicity	EducationLevel	BMI	Smoking	AlcoholConsumption	PhysicalActivity	DietQuality	...	Memory
0	1	67	0	3	0	37.205177	0	12.215677	7.780544	6.433890	...	
1	2	65	1	0	0	35.141843	1	17.111404	6.645284	1.112379	...	
2	3	62	0	1	1	17.875103	0	13.525546	9.585769	4.266008	...	
3	4	67	0	0	1	37.503437	1	19.952014	1.953946	6.797333	...	
4	5	65	1	0	2	29.187863	1	0.533209	8.759570	6.364302	...	

5 rows × 35 columns



Exploratory Data Analysis (EDA)

```
In [ ]: #Checking for NULL values
df.isna().sum()
```

Out[]:

	0
PatientID	0
Age	0
Gender	0
Ethnicity	0
EducationLevel	0
BMI	0
Smoking	0
AlcoholConsumption	0
PhysicalActivity	0
DietQuality	0
SleepQuality	0
FamilyHistoryAlzheimers	0
CardiovascularDisease	0
Diabetes	0
Depression	0
HeadInjury	0
Hypertension	0
SystolicBP	0
DiastolicBP	0
CholesterolTotal	0
CholesterolLDL	0
CholesterolHDL	0
CholesterolTriglycerides	0
MMSE	0
FunctionalAssessment	0
MemoryComplaints	0
BehavioralProblems	0
ADL	0
Confusion	0
Disorientation	0
PersonalityChanges	0
DifficultyCompletingTasks	0
Forgetfulness	0
Diagnosis	0
DoctorInCharge	0

dtype: int64

In []:

```
#Checking for duplicates
duplicates = df[df[list(df.columns)].duplicated()].index
duplicates
```

Out[]: Index([], dtype='int64')

The dataset got no duplicates and no null values

In []:

```
def reset_df():
    # 0.8-0.2 Train-val split
    df_train, df_val = train_test_split(df, test_size=0.2, random_state=5, stratify=df['Diagnosis'])

    train_diagnosis = df_train['Diagnosis'].reset_index(drop = True)
    val_diagnosis = df_val['Diagnosis'].reset_index(drop = True)


    #Dropping Irrelevant columns
    df_train = df_train.drop(['PatientID', 'DoctorInCharge', 'Diagnosis'], axis=1)
    df_val = df_val.drop(['PatientID', 'DoctorInCharge', 'Diagnosis'], axis=1)
```

```
#Standardizing every column except diagnosis
scaler = StandardScaler()

df_train = pd.DataFrame(scaler.fit_transform(df_train), columns=df_train.columns)
df_val = pd.DataFrame(scaler.transform(df_val), columns = df_val.columns)

#Adding diagnosis column which is not standardized
df_train['Diagnosis'] = train_diagnosis
df_val['Diagnosis'] = val_diagnosis
return df_train, df_val, scaler
df_train, df_val, scaler = reset_df()
```

```
In [ ]: #Plotting Correlation Matrix
correlation_matrix = df_train.corr()
plt.figure(figsize=(20, 16))
sns.heatmap(correlation_matrix, annot=False, cmap='coolwarm', center=0,
            square=True, linewidths=0.5, cbar_kws={"shrink": .8})
plt.title('Correlation Heatmap', fontsize=16)
plt.tight_layout()
plt.show()
```



```
In [ ]: def top_correlated_features(data, target_feature, top_n=10):
    # Identifies the top correlated features to a given target feature.

    # Calculate the correlation matrix
    correlation_matrix = data.corr()

    # Extract correlation values for the target feature
    target_correlation = correlation_matrix[target_feature].drop(labels=[target_feature])

    # Get the top N correlated features (absolute value to account for both positive and negative correlations)
    top_correlations = target_correlation.abs().sort_values(ascending=False).head(top_n)

    # Return the correlation values with original signs
    return target_correlation[top_correlations.index]
top_correlated_features(df_train,"BehavioralProblems")
```



Out []:

BehavioralProblems	
Diagnosis	0.239710
HeadInjury	0.064203
PersonalityChanges	-0.060831
Forgetfulness	0.060083
Smoking	-0.058980
Age	0.054790
FamilyHistoryAlzheimers	-0.050267
CholesterolHDL	0.047882
CardiovascularDisease	-0.044209
DiastolicBP	0.041060

dtype: float64

In []:


```
def plot_pdf(df, col):  
    #Plot two pdf: pdf of specific column given diagnosis = 0, pdf of specific column given diagnosis = 1  
    sns.set(style="whitegrid")  
  
    # Plot the PDF using Seaborn's kdeplot  
    plt.figure(figsize=(10, 6))  
    sns.histplot(df[df['Diagnosis'] == 0][col], kde=True, color='blue', label='Benign', stat="density", bins=10, alpha=0.5)  
    sns.histplot(df[df['Diagnosis'] == 1][col], kde=True, color='orange', label='Malignant', stat="density", bins=10, alpha=0.5)  
  
    plt.title(f"PDF of {col}", fontsize=16)  
    plt.xlabel(f"{col}", fontsize=14)  
    plt.ylabel("Density", fontsize=14)  
    plt.legend(title="Status")  
    plt.grid(True)  
  
    plt.show()  
plot_pdf(df_train,"FunctionalAssessment")
```



- 1. FunctionalAssessment > 0, tend to be Benign FunctionalAssessment < 0, tend to be Malignant
- 2. Graph overlaps, so additional features are needed to classify Benign/Malignant

In []:

```
plot_pdf(df_train,'MMSE')
```



```
In [ ]: def plot_box(df,col):
    #Plotting box plot of specific column for diagnosis = 0, diagnosis = 1

    plt.figure(figsize=(10, 6))
    sns.boxplot(data=df, x='Diagnosis', y=col, palette={'0': 'blue', '1': 'orange'})


    plt.title(f'Box Plot of {col} by Diagnosis', fontsize=16)
    plt.xlabel('Diagnosis', fontsize=14)
    plt.ylabel(col, fontsize=14)

    plt.grid(axis='y', linestyle='--', alpha=0.7)
    plt.show()
plot_box(df_train,"FunctionalAssessment")
```

<ipython-input-11-4e68a1b26e33>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df, x='Diagnosis', y=col, palette={'0': 'blue', '1': 'orange'})
```

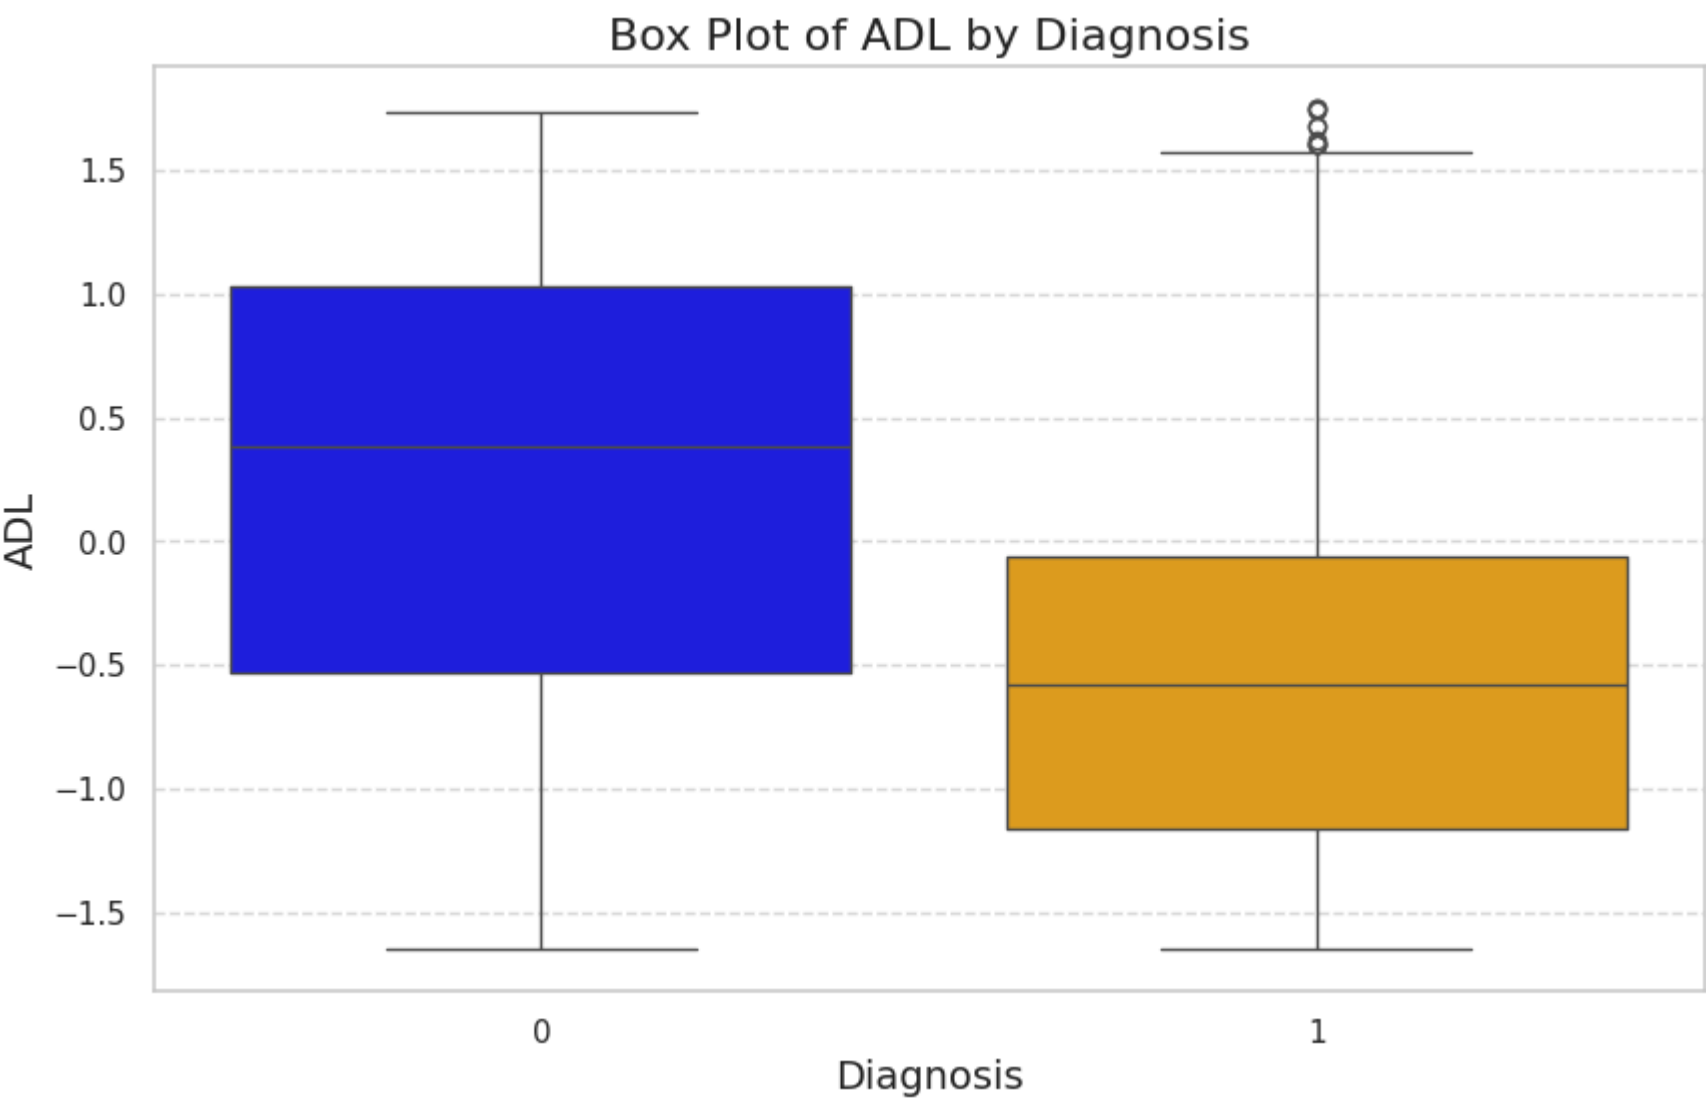



```
In [ ]: plot_box(df_train,"ADL")
```

<ipython-input-11-4e68a1b26e33>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df, x='Diagnosis', y=col, palette={'0': 'blue', '1': 'orange'})
```



```
In [ ]: #Finding upper bound for group 1
grouped = df_train.groupby('Diagnosis')['ADL']

# For Diagnosis = 1 (orange box)
Q1 = grouped.quantile(0.25).loc[1]
Q3 = grouped.quantile(0.75).loc[1]
IQR = Q3 - Q1
upper_bound = Q3 + 1.5 * IQR

# Print results for Diagnosis = 1
print("For Diagnosis = 1:")
print("Q1:", Q1)
print("Q3:", Q3)
print("IQR:", IQR)
print("Upper Bound:", upper_bound)
```

For Diagnosis = 1:
Q1: -1.1616764567839484
Q3: -0.059578148261666145
IQR: 1.1020983085222822
Upper Bound: 1.593569314521757

```
In [ ]: len(df_train[(df_train['ADL'] > upper_bound) & (df_train['Diagnosis'] == 1)])
```

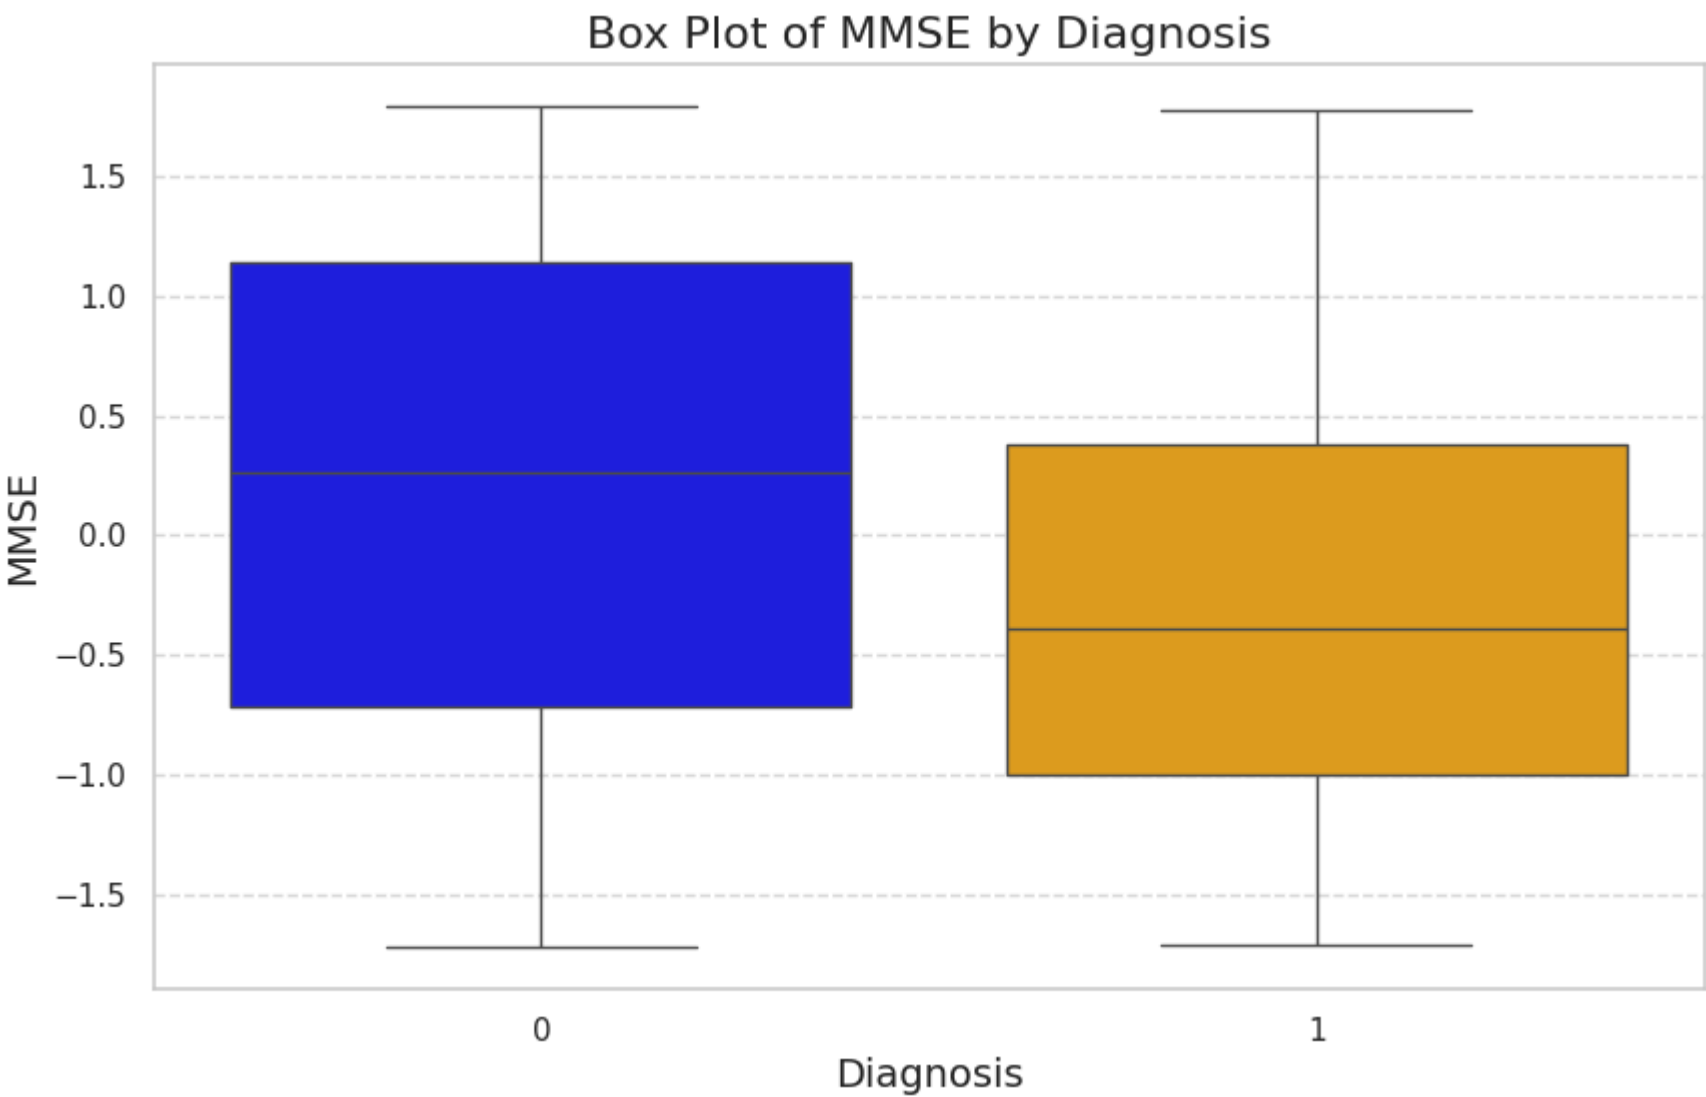
Out[]: 9

```
In [ ]: plot_box(df_train,"MMSE")
```

<ipython-input-11-4e68a1b26e33>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df, x='Diagnosis', y=col, palette={'0': 'blue', '1': 'orange'})
```

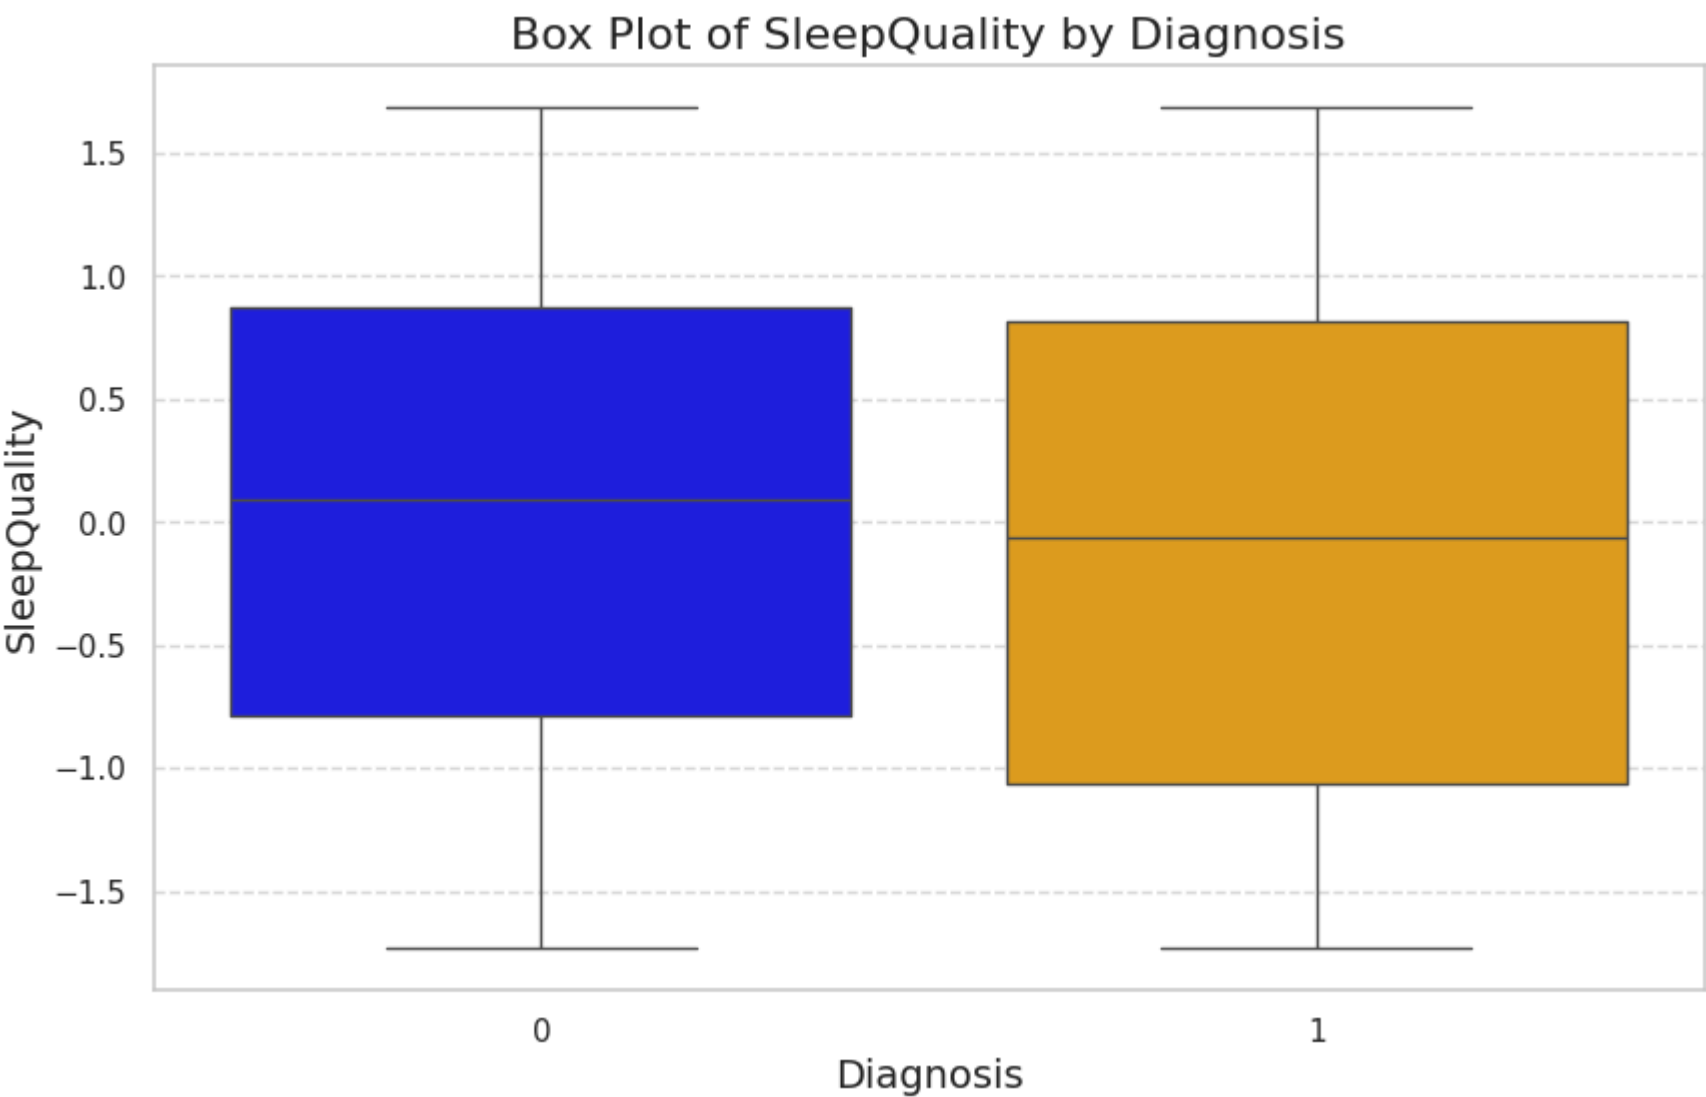


```
In [ ]: plot_box(df_train,"SleepQuality")
```

<ipython-input-11-4e68a1b26e33>:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df, x='Diagnosis', y=col, palette={'0': 'blue', '1': 'orange'})
```



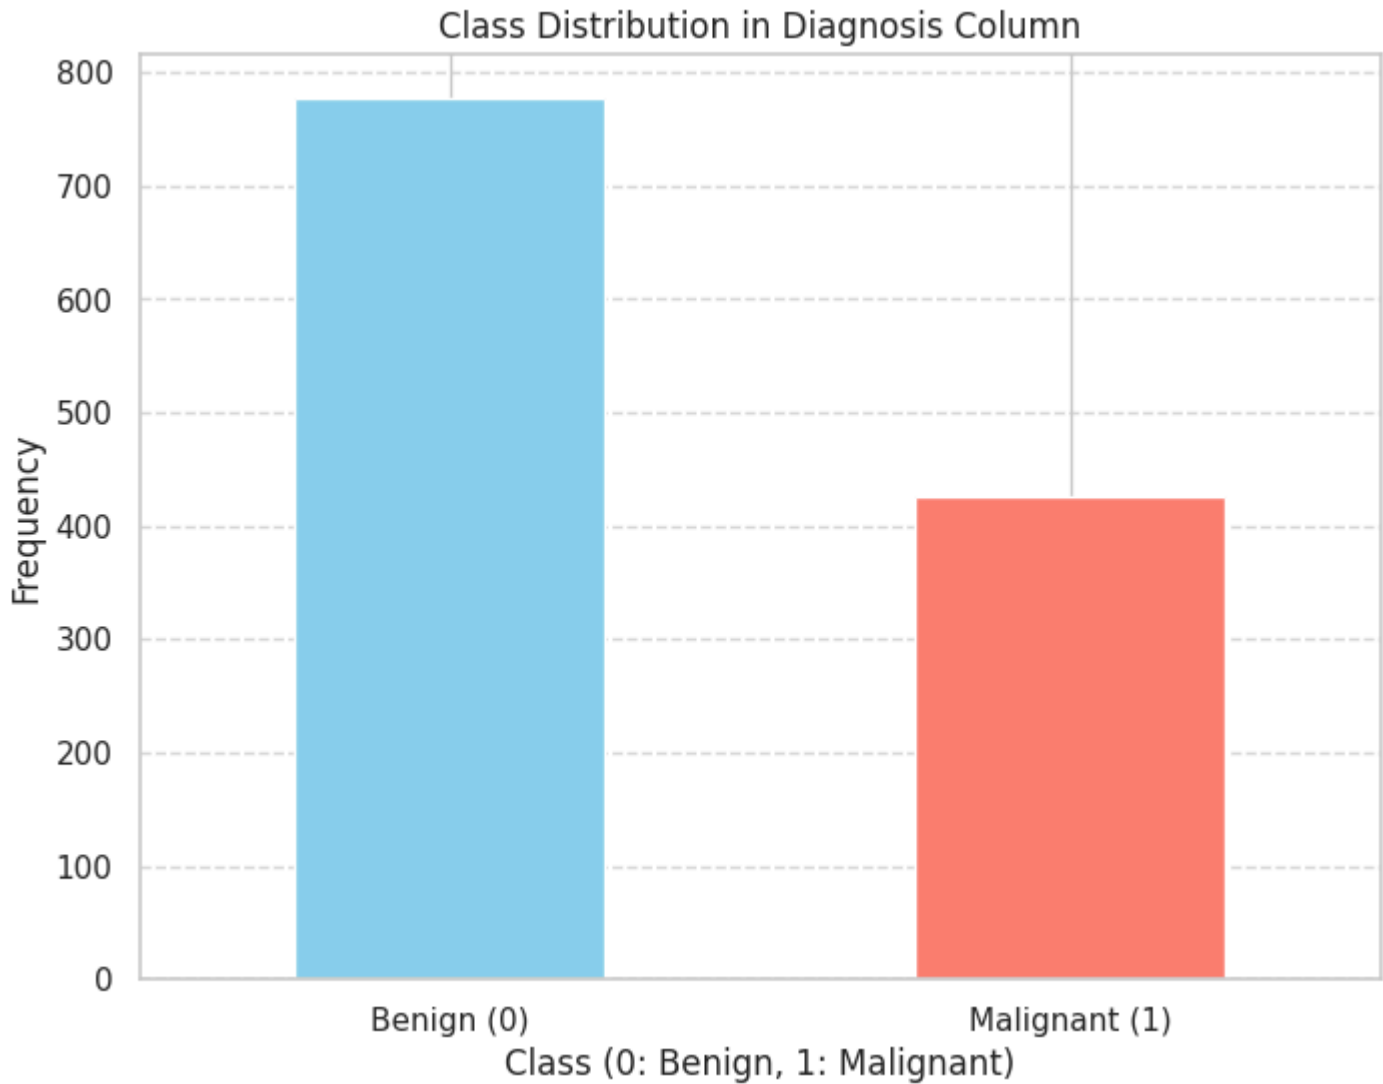
```
In [ ]: class_counts = df_train['Diagnosis'].value_counts()

# Print the distribution
print("Class Distribution:")
print(class_counts)

# Plot the distribution
plt.figure(figsize=(8, 6))
class_counts.plot(kind='bar', color=['skyblue', 'salmon'])
plt.title('Class Distribution in Diagnosis Column')
plt.xlabel('Class (0: Benign, 1: Malignant)')
```

```
plt.ylabel('Frequency')
plt.xticks([0, 1], ['Benign (0)', 'Malignant (1)'], rotation=0)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```

Class Distribution:
Diagnosis
0 777
1 426
Name: count, dtype: int64



Model

```
In [ ]: def evaluate(model, X, y, class_report = True):
        y_pred = model.predict(X)
        accuracy = accuracy_score(y,y_pred)
        f1 = f1_score(y,y_pred, average='weighted')
        sensitivity = recall_score(y,y_pred,average='weighted')
        balanced_accuracy = balanced_accuracy_score(y,y_pred)
        if class_report:
            print(classification_report(y, y_pred, target_names=["Benign","Malignant"]))
        return accuracy, f1, sensitivity, balanced_accuracy
```

```
In [ ]: def evaluate_model(model, X_train, y_train, X_val, y_val, class_report = True):
        train_accuracy, train_f1, train_sensitivity, train_balanced_accuracy = evaluate(model, X_train, y_train, class_report)
        validation_accuracy, validation_f1, validation_sensitivity, validation_balanced_accuracy = evaluate(model, X_val, y_val, c
        print(f"Training Metrics: Accuracy = {train_accuracy}, F1 score = {train_f1}, Sensitivity = {train_sensitivity}")
        print(f"Validation Metrics: Accuracy = {validation_accuracy}, F1 score = {validation_f1}, Sensitivity = {validation_sensit
```

```
In [ ]: def prepare_X_y(X_col, resample = True):
        #prepare X, y given arbitrary training features with SMOTE resampling
        X_train = df_train[X_col]
        y_train = df_train['Diagnosis']

        X_val = df_val[X_col]
        y_val = df_val['Diagnosis']

        #Apply SMOTE resampling
        if resample == True:
            smote = SMOTE(random_state=0)
            X_train, y_train = smote.fit_resample(X_train, y_train)
        return X_train, y_train, X_val, y_val
```

Logistic Regression

```
In [ ]: #Try using MMSE, FunctionalAssessment, MemoryComplaints, BehavioralProblems, ADL for fitting (5 most important features)
X_train, y_train, X_val, y_val = prepare_X_y(['MMSE', 'FunctionalAssessment', 'MemoryComplaints', 'BehavioralProblems', 'ADL',
lr_model = LogisticRegression(
    penalty='l2', # L2 regularization (ridge)
    C=1.0,
    solver='liblinear', # Good for small datasets
    random_state=42,
    max_iter=1000,
```

```
)  
  
# Fit the model  
lr_model.fit(X_train, y_train)  
evaluate_model(lr_model, X_train, y_train, X_val, y_val)
```

	precision	recall	f1-score	support
Benign	0.84	0.82	0.83	777
Malignant	0.82	0.84	0.83	777
accuracy			0.83	1554
macro avg	0.83	0.83	0.83	1554
weighted avg	0.83	0.83	0.83	1554

	precision	recall	f1-score	support
Benign	0.88	0.80	0.84	195
Malignant	0.69	0.80	0.74	106
accuracy			0.80	301
macro avg	0.78	0.80	0.79	301
weighted avg	0.81	0.80	0.80	301

Training Metrics: Accuracy = 0.832046332046332, F1 score = 0.8320262301382366, Sensitivity = 0.832046332046332
Validation Metrics: Accuracy = 0.8006644518272426, F1 score = 0.803641904264886, Sensitivity = 0.8006644518272426

Decision Tree


```
In [ ]: #Decision Tree  
clf = DecisionTreeClassifier(random_state=0, max_depth = 5, min_samples_split = 2)  
clf.fit(X_train, y_train)  
evaluate_model(clf, X_train, y_train, X_val, y_val)
```

	precision	recall	f1-score	support
Benign	0.95	0.98	0.97	777
Malignant	0.98	0.95	0.96	777
accuracy			0.96	1554
macro avg	0.96	0.96	0.96	1554
weighted avg	0.96	0.96	0.96	1554

	precision	recall	f1-score	support
Benign	0.96	0.92	0.94	195
Malignant	0.87	0.93	0.90	106
accuracy			0.93	301
macro avg	0.92	0.93	0.92	301
weighted avg	0.93	0.93	0.93	301

Training Metrics: Accuracy = 0.9646074646074646, F1 score = 0.9646010002174386, Sensitivity = 0.9646074646074646
Validation Metrics: Accuracy = 0.9269102990033222, F1 score = 0.927473865474596, Sensitivity = 0.9269102990033222

```
In [ ]: import shap  
  
# Initialize SHAP explainer  
explainer = shap.TreeExplainer(clf)  
  
# Get SHAP values for the entire dataset  
shap_values = explainer.shap_values(X_val)  
  
# Visualize SHAP values  
shap.summary_plot(shap_values[:, :, 1], X_val) # For binary classification (class 1)
```



Random Forest

```
In [ ]: #Random Forest using all features & Grid search
features = [col for col in df_train.columns if col != 'Diagnosis']
X_train, y_train, X_val, y_val = prepare_X_y(features)
```

```
In [ ]: param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [None, 10],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2, 3],
    'max_features': ['log2', 'sqrt'],
    'bootstrap': [True]
}
```

```
In [ ]: rf = RandomForestClassifier(random_state=42)
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, scoring='precision', verbose=1) #Accuracy
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 48 candidates, totalling 240 fits

/usr/local/lib/python3.10/dist-packages/numpy/ma/core.py:2820: RuntimeWarning: invalid value encountered in cast
_data = np.array(data, dtype=dtype, copy=copy,

```
Out [ ]: GridSearchCV
  ▸ best_estimator_: RandomForestClassifier
    ▸ RandomForestClassifier
```

```
In [ ]: print("Best Parameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
best_rf = grid_search.best_estimator_
```

Best Parameters: {'bootstrap': True, 'max_depth': None, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}


Best Accuracy: 0.9760440649975571

```
In [ ]: #Feature importances from our Random Forest
importances = best_rf.feature_importances_

feature_names = X_train.columns

indices = np.argsort(importances)[::-1]

# Plot the feature importances
plt.figure(figsize=(10, 6))
plt.title("Feature Importances")
plt.barh(range(len(importances)), importances[indices], align="center")
plt.yticks(range(len(importances)), np.array(feature_names)[indices])
plt.xlabel("Relative Importance")
plt.show()
```



```
In [ ]: validation_accuracy, validation_f1, validation_sensitivity, validation_balanced_accuracy = evaluate(best_rf, X_val, y_val)
print(f"Validation Metrics: Accuracy = {validation_accuracy}, F1 score = {validation_f1}, Sensitivity = {validation_sensitivity}")
```

	precision	recall	f1-score	support
Benign	0.93	0.95	0.94	195
Malignant	0.90	0.87	0.88	106
accuracy			0.92	301
macro avg	0.92	0.91	0.91	301
weighted avg	0.92	0.92	0.92	301

Validation Metrics: Accuracy = 0.920265780730897, F1 score = 0.9199038477754095, Sensitivity = 0.920265780730897

XGBoost

```
In [ ]: #Fitting XGBoost on 30 most important features
X_train, y_train, X_val, y_val = prepare_X_y(np.array(feature_names)[indices][:30])

xgb_clf = xgb.XGBClassifier(
    objective='binary:logistic',
    random_state=42,
    n_jobs=-1, # Use all available cores
)

xgb_clf.fit(X_train,y_train)
evaluate_model(xgb_clf,X_train, y_train, X_val, y_val)
```

	precision	recall	f1-score	support
Benign	1.00	1.00	1.00	777
Malignant	1.00	1.00	1.00	777
accuracy			1.00	1554
macro avg	1.00	1.00	1.00	1554
weighted avg	1.00	1.00	1.00	1554

	precision	recall	f1-score	support
Benign	0.97	0.93	0.95	195
Malignant	0.88	0.94	0.91	106
accuracy			0.93	301
macro avg	0.92	0.94	0.93	301
weighted avg	0.94	0.93	0.93	301

Training Metrics: Accuracy = 1.0, F1 score = 1.0, Sensitivity = 1.0
Validation Metrics: Accuracy = 0.9335548172757475, F1 score = 0.9340671504314509, Sensitivity = 0.9335548172757475

```
In [ ]: #XGBoost with grid search
param_grid = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 1],
    'max_depth': [3, 5, 7],
}

# Create the XGBoost classifier
xgb_clf = xgb.XGBClassifier(
    objective='binary:logistic',
    random_state=42,
    n_jobs=-1 # Use all available cores
)

grid_search = GridSearchCV(estimator=xgb_clf, param_grid=param_grid, cv=5, scoring='precision', verbose=1)

grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

Out []:

GridSearchCV ⓘ ?

▶ best_estimator_: XGBClassifier

▶ XGBClassifier

```
In [ ]: print("Best Parameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
best_xgb = grid_search.best_estimator_
```

Best Parameters: {'learning_rate': 0.01, 'max_depth': 5, 'n_estimators': 100}
Best Accuracy: 0.9735702856249894

```
In [ ]: validation_accuracy, validation_f1, validation_sensitivity, validation_balanced_accuracy = evaluate(best_xgb, X_val, y_val)
print(f"Validation Metrics: Accuracy = {validation_accuracy}, F1 score = {validation_f1}, Sensitivity = {validation_sensitivity}")
```


	precision	recall	f1-score	support
Benign	0.96	0.95	0.95	195
Malignant	0.91	0.92	0.92	106
accuracy			0.94	301
macro avg	0.93	0.94	0.93	301
weighted avg	0.94	0.94	0.94	301

Validation Metrics: Accuracy = 0.9401993355481728, F1 score = 0.9403246524815795, Sensitivity = 0.9401993355481728

```
In [ ]: #SKlearn gradient boosting + Grid search
from sklearn.ensemble import GradientBoostingClassifier
model = GradientBoostingClassifier(random_state = 0)

grid_search = GridSearchCV(model, param_grid, cv = 5, scoring = 'precision')
grid_search.fit(X_train,y_train)
best_gb = grid_search.best_estimator_
```

```
In [ ]: print("Best Parameters:", grid_search.best_params_)
print("Best Accuracy:", grid_search.best_score_)
```

Best Parameters: {'learning_rate': 0.01, 'max_depth': 5, 'n_estimators': 50}
Best Accuracy: 0.9643439784819096

```
In [ ]: validation_accuracy, validation_f1, validation_sensitivity, validation_balanced_accuracy = evaluate(best_gb, X_val, y_val)
print(f"Validation Metrics: Accuracy = {validation_accuracy}, F1 score = {validation_f1}, Sensitivity = {validation_sensitivity}")
```

	precision	recall	f1-score	support
Benign	0.96	0.93	0.94	195
Malignant	0.88	0.92	0.90	106
accuracy			0.93	301
macro avg	0.92	0.93	0.92	301
weighted avg	0.93	0.93	0.93	301

Validation Metrics: Accuracy = 0.9269102990033222, F1 score = 0.9273451072876345, Sensitivity = 0.9269102990033222

```
In [ ]: #CatBoost + Gridsearch
from catboost import CatBoostClassifier
cbc = CatBoostClassifier(verbose=0)
grid_search = GridSearchCV(estimator=cbc, param_grid=param_grid, cv=5, scoring='precision', verbose=1)

grid_search.fit(X_train,y_train)
best_cbc = grid_search.best_estimator_
```

Fitting 5 folds for each of 27 candidates, totalling 135 fits

Gridsearch Visualization (Edited)


```
In [ ]: # taking results
results = grid_search.cv_results_

# checking result tables
results_df = pd.DataFrame(results)
print(results_df[['param_learning_rate', 'param_max_depth', 'param_n_estimators', 'mean_test_score']])

# compute score
scores = results['mean_test_score']
iterations = range(1, len(scores) + 1)

# Line graph
plt.plot(iterations, scores, marker='o', label='Validation Score')
plt.xlabel('Iteration')
plt.ylabel('Mean Test Score')
plt.title('GridSearchCV Progress')
plt.grid()
plt.legend()
plt.show()
```

	param_learning_rate	param_max_depth	param_n_estimators	mean_test_score
0	0.01	3	50	0.965752
1	0.01	3	100	0.968102
2	0.01	3	200	0.969561
3	0.01	5	50	0.970811
4	0.01	5	100	0.971982
5	0.01	5	200	0.973323
6	0.01	7	50	0.969303
7	0.01	7	100	0.968049
8	0.01	7	200	0.970737
9	0.10	3	50	0.966932
10	0.10	3	100	0.968125
11	0.10	3	200	0.963040
12	0.10	5	50	0.965520
13	0.10	5	100	0.968360
14	0.10	5	200	0.964393
15	0.10	7	50	0.968784
16	0.10	7	100	0.968724
17	0.10	7	200	0.970035
18	1.00	3	50	0.931027
19	1.00	3	100	0.932732
20	1.00	3	200	0.932717
21	1.00	5	50	0.949115
22	1.00	5	100	0.956269
23	1.00	5	200	0.951239
24	1.00	7	50	0.943992
25	1.00	7	100	0.951719
26	1.00	7	200	0.951717



Learning curve (Edited)

```
In [ ]: from sklearn.model_selection import learning_curve
# define model
model = CatBoostClassifier(**best_cbc.get_params())

# calculating learning curve
train_sizes, train_scores, test_scores = learning_curve(
    model, X_train, y_train, cv=10, scoring='accuracy', n_jobs=-1, train_sizes=np.linspace(0.1, 1.0, 10)
)


train_mean = np.mean(train_scores, axis=1)
train_std = np.std(train_scores, axis=1)
test_mean = np.mean(test_scores, axis=1)
test_std = np.std(test_scores, axis=1)

# plot
plt.figure(figsize=(10, 6))
plt.plot(train_sizes, train_mean, 'o-', color='blue', label='Training Score')
plt.plot(train_sizes, test_mean, 'o-', color='orange', label='Validation Score')

# std area
plt.fill_between(train_sizes, train_mean - train_std, train_mean + train_std, alpha=0.2, color='blue')
plt.fill_between(train_sizes, test_mean - test_std, test_mean + test_std, alpha=0.2, color='orange')

plt.title('Learning Curve')
plt.xlabel('Training Examples')
plt.ylabel('Accuracy')
plt.legend(loc='best')
```

```
plt.grid()  
plt.show()
```



```
In [ ]: print("Best Parameters:", grid_search.best_params_)  
        print("Best Accuracy:", grid_search.best_score_)
```

Best Parameters: {'learning_rate': 0.01, 'max_depth': 5, 'n_estimators': 200}
Best Accuracy: 0.9733232763028855

```
In [ ]: validation_accuracy, validation_f1, validation_sensitivity, validation_balanced_accuracy = evaluate(best_cbc, X_val, y_val)  
        print(f"Validation Metrics: Accuracy = {validation_accuracy}, F1 score = {validation_f1}, Sensitivity = {validation_sensitivity}")
```



	precision	recall	f1-score	support
Benign	0.97	0.95	0.96	195
Malignant	0.91	0.94	0.93	106
accuracy			0.95	301
macro avg	0.94	0.95	0.94	301
weighted avg	0.95	0.95	0.95	301

Validation Metrics: Accuracy = 0.946843853820598, F1 score = 0.9470606199127708, Sensitivity = 0.946843853820598

Confusion Matrix and AUC (Edited)

```
In [ ]: #added for confusion matrix  
        from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay  
        import matplotlib.pyplot as plt  
  
        y_true = y_val  
        y_pred = best_cbc.predict(X_val)  
  
        #compute  
        cm = confusion_matrix(y_true, y_pred)  
  
        #display  
        plt.figure(figsize=(8, 6))  
        sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')  
        plt.xlabel('Predicted Labels')  
        plt.ylabel('True Labels')  
        plt.title('Confusion Matrix')  
        plt.show()  
  
        #AUC  
        import sklearn.metrics as metrics  
        probs = best_cbc.predict_proba(X_val)  
        preds = probs[:,1]  
        fpr, tpr, threshold = metrics.roc_curve(y_true, preds)  
        roc_auc = metrics.auc(fpr, tpr)  
  
        #method: plt  
        import matplotlib.pyplot as plt  
        plt.title('Receiver Operating Characteristic')  
        plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc)
```

```
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



In []: *#Refitting the model on whole data*

```
#Top 30 features
features = list(X_train.columns)

diag = df['Diagnosis'].reset_index(drop = True)

#Drop Irrelevant Columns
df_scaled = df.drop(['PatientID', 'DoctorInCharge', 'Diagnosis'], axis=1)

#Standardizing every column except diagnosis
df_scaled = pd.DataFrame(scaler.transform(df_scaled), columns=df_scaled.columns)

#Adding unscaled diagnosis column back
df_scaled['Diagnosis'] = diag

#Get X,y
X = df_scaled[features]
```

```
y = df_scaled['Diagnosis']
smote = SMOTE(random_state=20)

#Resample
X, y = smote.fit_resample(X, y)

#Refit the model on whole data
best_model = CatBoostClassifier(**best_cbc.get_params())
best_model.fit(X, y)
```

Out[]: <catboost.core.CatBoostClassifier at 0x7bedacefca30>

Ensembling

```
In [ ]: X_train, y_train, X_val, y_val = prepare_X_y(['MMSE', 'FunctionalAssessment', 'MemoryComplaints', 'BehavioralProblems', 'ADL'],

#Logistic regression
lr_pred = lr_model.predict_proba(X_val)

#Decision Tree
dt_pred = clf.predict_proba(X_val)

#Random Forest
features = [col for col in df_train.columns if col != 'Diagnosis']
X_train, y_train, X_val, y_val = prepare_X_y(features)

rf_pred = best_rf.predict_proba(X_val)

#Tree Boosting
X_train, y_train, X_val, y_val = prepare_X_y(np.array(feature_names)[indices][:30])
xgb_pred = best_xgb.predict_proba(X_val)
gb_pred = best_gb.predict_proba(X_val)
cbc_pred = best_cbc.predict_proba(X_val)
```

```
In [ ]: #Weighted Averaging
y_pred_prob = (0.05*lr_pred + 0.05*dt_pred + 0.2*rf_pred + 0.2*xgb_pred + 0.2*gb_pred + 0.3*cbc_pred)
y_pred = np.argmax(y_pred_prob, axis=1)
print(f"Accuracy: {accuracy_score(y_val,y_pred)}")
print(f"Recall: {recall_score(y_val,y_pred, average='weighted')}")
```

Accuracy: 0.9435215946843853
Recall: 0.9435215946843853

Model Submission

```
In [ ]: df_bf = pd.read_csv('best_submission.csv')
df_test = pd.read_csv("test.csv")
```

```
In [ ]: def get_submission(model, features_used, scaler):
    df_test = pd.read_csv("test.csv")
    pat_ID = pd.DataFrame(df_test['PatientID'])
    df_test = df_test.drop(['PatientID', 'DoctorInCharge'], axis=1)
    df_test = pd.DataFrame(scaler.transform(df_test), columns=df_test.columns)

    X_t = df_test[features_used]
    y_test_pred = model.predict(X_t)
    pat_ID['Diagnosis'] = y_test_pred
    pat_ID.to_csv('submission.csv', index=False)

    accuracy, f1, sensitivity, _ = evaluate(model, X_t, df_bf['Diagnosis'])
    print(f"Comparing best result: Accuracy = {accuracy}, F1 score = {f1}, Sensitivity = {sensitivity}")

    print("Submission File Downloaded")
```

```
In [ ]: get_submission(best_model, list(X_train.columns), scaler)
```

	precision	recall	f1-score	support
Benign	1.00	1.00	1.00	426
Malignant	1.00	1.00	1.00	219
accuracy			1.00	645
macro avg	1.00	1.00	1.00	645
weighted avg	1.00	1.00	1.00	645

Comparing best result: Accuracy = 1.0, F1 score = 1.0, Sensitivity = 1.0
Submission File Downloaded

```
In [ ]: %%shell
jupyter nbconvert --to html /content/Colab_downladed_file_name.ipynb
```